

# Assignment #1: Linux

**Due Date 1:** Friday, January 22, 2021, 5:00 pm EST

**Due Date 2:** Friday, January 29, 2021, 5:00 pm EST

**Online Quiz:** Wednesday, January 27, 2021, 5:00 pm EST

Questions 1 and 2 are due on Due Date 1; the remainder of the assignment is due on Due Date 2. If you joined the course within two days of Due Date 1, then the entire assignment is due on Due Date 2. You must submit the online quiz on Learn by the Quiz date.

Topics that must have been completed before starting Due Date 1:

1. Linux: The Teaching Environment
2. Linux: Interacting with the Shell
3. Linux: Directories and Files
4. Linux: Regular Expressions
5. Handouts: Getting Started
6. Handouts: Linux Commands

Topics that must have been completed before starting Due Date 2:

1. Linux: bash scripts
2. Software Testing

**Note:** We suggest creating the directory `~/cs246/w21/a1` and creating all the assignment solutions in this directory.

## bash and Regular Expressions

1. **10 marks.** Provide a Linux command line to accomplish each of the following tasks. Your answer in each subquestion should consist of a single command or pipeline of commands, with no separating semicolons (;). (Please verify before submitting that your solution consists of a single line. Use `wc` for this.) Before beginning this question, familiarize yourself with the commands outlined on the Linux handout. Keep in mind that some commands have options not listed on the sheet, so you may need to examine some man pages. Note that some tasks refer to a file `myfile.txt`. No `myfile.txt` is given. You should create your own for testing.
  - (a) Print the 10th through 25th words (including the 10th and 25th words) in `/usr/share/dict/words`. You may take advantage of the fact that the words in this file are each on a separate line. Place your command pipeline in the file `alq1a.txt`.
  - (b) Print the (non-hidden) contents of the current directory in reverse of the normal order. Place your command pipeline in the file `alq1b.txt`.
  - (c) Print the number of lines in the text file `myfile.txt` that do *not* contain the string `cs246` (all in lower-case). Place your command pipeline in the file `alq1c.txt`.
  - (d) Print the first line that contains the string `cs246` (all in lower-case) from the text file `myfile.txt`. Place your command pipeline in the file `alq1d.txt`.
  - (e) Print the number of lines in the text file `myfile.txt` that contain the string `linux.student.cs.uwaterloo.ca` where each letter could be either upper-case or lower-case. (Hint: this is not as obvious as you may think—carefully re-read the special symbols and their meanings first!) Place your command pipeline in the file `alq1e.txt`.

- (f) Print all (non-hidden) files/directories in any *subdirectory* of the current directory that end with lower-case `.c` (immediate subdirectories only, not subdirectories of subdirectories). Do not use `find`. (Hint: there's an easy way to do this using only `ls` if you're creative with globbing patterns.) Place your command pipeline in the file `alq1f.txt`.
- (g) Out of the first 20 lines of `myfile.txt`, how many contain at least one digit? Place the command pipeline that prints this number in the file `alq1g.txt`.
- (h) Print all (non-hidden) files in the current directory that start with `a`, contain at least one `b`, and end with `.c` (these required letters must be in lower-case in order to match). Place your command pipeline in the file `alq1h.txt`.
- (i) Print a listing, in long form, of all non-hidden entries (files, directories, etc.) in the current directory that are executable by at least one of owner, group, other (the other permission bits could be anything). Do not attempt to solve this problem with `find`. Place your command pipeline in the file `alq1i.txt`.
- (j) Before attempting this subquestion, do some reading (either skim the man page or have a look on the Web) on the `awk` utility. In particular, be sure you understand the effect of the command

```
awk '{print $1}' < myfile.txt
```

Give a Linux pipeline that gives a sorted, duplicate-free list of userids currently signed on to the (school) machine the command is running on.

Place your command pipeline in the file `alq1j.txt`.

2. **10 marks.** For each of the following text search criteria, provide a regular expression that matches the criterion, suitable for use with `egrep`. Your answer in each case should be a text file that contains just the regular expression (i.e., you don't need to include the `egrep` command in your submitted solution), on a single line (again, use `wc` to verify this). If your pattern contains special characters, enclose it in quotes.

- (a) Lines that contain both `cs246` and `cs247`, in lower-case.  
Place your answer in the file `alq2a.txt`.
- (b) Lines that contain nothing but a single occurrence of laughter, where laughter is defined as a string of the form `Hahahahahahahahahahahahaha!`, with arbitrarily many `ha`'s. The string must start with `H` and end with `!`. Place your answer in the file `alq2b.txt`.
- (c) Lines that contain nothing but a single occurrence of generalized laughter, which is like ordinary laughter, except that there can be arbitrarily many (but at least one) `a`'s between each pair of consecutive `h`'s. (For example: `Haahahaaaa!`) Place your answer in the file `alq2c.txt`.
- (d) Lines that contain at least one lower-case `a` and at least two lower-case `b`'s.  
Place your answer in the file `alq2d.txt`.
- (e) Lines consisting of a definition of a single C variable of type `int`, without initialization, optionally preceded by `unsigned`, and optionally followed by any single line `// comment`. Example:

```
int varname; // optional comment
```

You may assume that all of the whitespace in the line consists of space characters (no tabs). You may also assume that `varname` will not be a C keyword (i.e., you do not have to try to check for this with your regular expression). If you don't remember the rules for naming a C variable, please consult <https://www.programiz.com/c-programming/c-variables-constants> Place your answer in the file `alq2e.txt`.

## Testing tools

**Note: the scripts you write in the following questions will be useful every time you write a program. Be sure to complete them!** In this course, you will be responsible for your own testing. As you fix bugs and refine your code, you will very often need to rerun old tests, to check that existing bugs have been fixed, and to ensure that no new bugs have been introduced. This task is *greatly* simplified if you take the time to create a formal test suite, and build tools to automate your testing. In the following questions, you will implement such tools as `bash` scripts.

3. **5 marks.** Write a `bash` script called `produceOutputs` that is invoked as follows:

```
./produceOutputs suite-file program
```

The argument `suite-file` is the name of a file containing a list of filename stems (more details below), and the argument `program` is the name of a program to be run.

The `produceOutputs` script runs `program` on each test in the test suite and, for each test, creates a file that contains the output produced for that test.

The file `suite-file` contains a list of stems, from which we construct the names of files containing the command-line arguments used by each test. Stems will not contain spaces. For example, suppose our suite file is called `suite.txt` and contains the following entries:

```
test1 test2
reallyBigTest
```

Then our test suite consists of three tests. The first one (`test1`) will use the file `test1.args`. The second one (`test2`) will use the file `test2.args`. The last one (`reallyBigTest`) will use the file `reallyBigTest.args`.

A sample run of `produceOutputs` would be as follows:

```
./produceOutputs suite.txt ./myprogram
```

The script will then run `./myprogram` three times, once for each test specified in `suite.txt`:

- The first time, it will run `./myprogram` with command-line arguments provided to the program from `test1.args`. The results, captured from standard output, will be stored in `test1.out`.
- The second time, it will run `./myprogram` with command-line arguments provided to the program from `test2.args`. The results, captured from standard output, will be stored in `test2.out`.
- The third time, it will run `./myprogram` with command-line arguments provided to the program from `reallyBigTest.args`. The results, captured from standard output, will be stored in `reallyBigTest.out`.

Note that if the test suite contains a stem but a corresponding `.args` file is not present, the program is run without providing any command-line arguments.

Your script must also check for incorrect number of command-line arguments to `produceOutputs`. If such an error condition arises, print an informative error message to standard error (the exact message is up to you) and abort the script with a **non-zero** exit status.

You can find an example of the expected output of `produceOutputs` in the directory `a1/q3` on your Git repository. Using the contents of this directory, running your script as:

```
./produceOutputs test_suite.txt ./my_factorial_correct
```

should produce files named `test1.out`, `test2.out`, and `test3.out` identical to those provided to you as examples in the same directory.

**Note on purpose of this script:** This script will be useful in situations where we provide you with a binary version of a program (but not its source code) that you must implement. By creating your own test cases (`.args` files) and then using this script to produce the intended output you will have something to compare with when you implement your own solution (see the next question for how to automate the comparisons).

4. **10 marks.** Create a bash script called `runSuite` that is invoked as follows:

```
./runSuite suite-file program
```

The argument `suite-file` is the name of a file containing a list of filename stems (more details below), and the argument `program` is the name of the program to be run.

In summary, the `runSuite` script runs `program` on each test in the test suite (as specified by `suite-file`) and reports on any tests whose output does not match the expected output.

The file `suite-file` contains a list of stems, from which we construct the names of files containing the command-line arguments and expected output of each test. Stems will not contain spaces. For example, suppose our suite file is called `suite.txt` and contains the following entries:

```
test1 test2
reallyBigTest
```

Then our test suite consists of three tests. The first one (`test1`) will use the file `test1.args` to hold its command-line arguments, and `test1.out` to hold its expected output. The second one (`test2`) will use the file `test2.args` to hold its command-line arguments, and `test2.out` to hold its expected output. The last one (`reallyBigTest`) will use the file `reallyBigTest.args` to hold its command-line arguments, and `reallyBigTest.out` to hold its expected output.

A sample run of `runSuite` would be as follows:

```
./runSuite suite.txt ./myprogram
```

The script will then run `./myprogram` three times, once for each test specified in `suite.txt`:

- The first time, it will run `./myprogram` with command-line arguments provided to the program from `test1.args`. The results, captured from standard output, will be compared with the contents of `test1.out`.
- The second time, it will run `./myprogram` with command-line arguments provided to the program from `test2.args`. The results, captured from standard output, will be compared with the contents of `test2.out`.
- The third time, it will run `./myprogram` with command-line arguments provided to the program from `reallyBigTest.args`. The results, captured from standard output, will be compared with the contents of `reallyBigTest.out`.

Note that if the test suite contains a stem but a corresponding `.args` file is not present, the program is run without providing any command-line arguments.

If the output of a given test case differs from the expected output, print the following to standard output (assuming test `test2` failed):

```
Test failed: test2
Args:
(contents of test2.args, if it exists)
Expected:
(contents of test2.out)
Actual:
(contents of the actual program output)
```

with the `(contents ...)` lines replaced with actual file contents, without any changes, as described. The literal output `Args:` must appear, even if the corresponding file does not exist. Note that there is no whitespace after the colon for each of `Args:`, `Expected:`, and `Actual:` except for the newline character.

**Follow these output specifications very carefully. You will lose a lot of marks if your output does not match them.** If you need to create temporary files, create them in `/tmp`, and use the `mktemp` command to prevent name duplications. **Also be sure to delete any temporary files you create in `/tmp`.**

You can find an example of the expected output of `runSuite` in the directory `a1/q4` on your Git repository. Using the contents of this directory, running your script as:

```
./runSuite test_suite.txt ./my_factorial_buggy
```

should produce an output identical to the example provided to you in the file `buggy_test_results.out` in the same directory.

**Note:** Do **NOT** attempt to compare outputs by storing them in shell variables, and then comparing the shell variables. This is a very bad idea, and it does not scale well to programs that produce large outputs. We reserve the right to deduct marks (on this and all assignments) for bad solutions like this would be.

**You can get most of the marks for this question by fulfilling the above requirements. For full marks, your script must also check for the following error conditions:**

- incorrect number of command-line arguments to `runSuite`
- missing or unreadable `.out` files (for example, the suite file contains an entry `xxx`, but `xxx.out` doesn't exist or is unreadable).

If such an error condition arises, print an informative error message to standard error and abort the script with a **non-zero** exit status.

5. **15 marks.** In this question, you will generalize the `produceOutputs` and `runSuite` scripts that you created in problems 3 and 4. As they are currently written, these scripts cannot be used with programs that take input from standard input. For this problem, you will enhance `produceOutputs` and `runSuite` so that, in addition to (optionally) passing command-line arguments to the program being executed, the program can also be (optionally) provided input from standard input. The interface to the scripts remains the same:

```
./produceOutputs suite.txt ./myprogram
./runSuite suite.txt ./myprogram
```

The format of the suite file remains the same. But now, for each `testname` in the suite file, there might be an optional `testname.in`. If the file `testname.in` is present, then the script (`produceOutputs` or `runSuite`) will run `myprogram` with the contents of `testname.args` passed on the command-line as before and the contents of `testname.in` used for input on `stdin`. If `testname.in` is not present, then the behaviour is almost identical to problem 3/4 (see below for a difference in the output): `myprogram` is run with command-line arguments coming from `testname.args` with nothing supplied as input on standard input.

The output of `runSuite` is changed to now also show the input provided to a test if the test failed. Assuming test `test2` from Q4 failed, the output generated by the updated `runSuite` is as follows:

```
Test failed: test2
Args:
  (contents of test2.args, if it exists)
Input:
  (contents of test2.in, if it exists)
Expected:
  (contents of test2.out)
Actual:
  (contents of the actual program output)
```

with the `(contents ...)` lines replaced with actual file contents, as described. The literal output `Args:` and `Input:` must appear, even if the corresponding files do not exist. Note that there is no whitespace after the colon for each of `Args:`, `Input:`, `Expected:`, and `Actual:` except for the newline character.

You can find an example of the expected output of the updated `produceOutputs` and `runSuite` in the directory `a1/q5` on your Git repository. Using the contents of this directory, running your script as:

```
./produceOutputs test_suite.txt ./my_factorial_correct
```

should produce files named `test1.out`, `test2.out`, and `test3.out` identical to those provided to you as examples in the same directory.

And using the contents of this directory, running your script as:

```
./runSuite test_suite.txt ./my_factorial_buggy
```

should produce an output identical to the example provided to you in the file `buggy_test_results.out` in the same directory.

**All error-checking that was required in problems 3 and 4 is required here as well.**

- (a) Modify `produceOutputs` to handle input from standard input.
- (b) Modify `runSuite` to handle input from standard input.

**Note:** To get this working should require only very small changes to your solution to problems 3 and 4.

## Submission

The following files are due at Due Date 1: `alq1a.txt`, `alq1b.txt`, `alq1c.txt`, `alq1d.txt`, `alq1e.txt`, `alq1f.txt`, `alq1g.txt`, `alq1h.txt`, `alq1i.txt`, `alq1j.txt`, `alq2a.txt`, `alq2b.txt`, `alq2c.txt`, `alq2d.txt`, `alq2e.txt`.  
The following files are due at Due Date 2: `produceOutputs`, `runSuite`, `produceOutputs`, `runSuite`.